

Round 2 Evaluation Metrics

BY,
Krithik S

I evaluated my selected models on 1,185 held-out posts after choosing models on a separate validation split. I remove duplicate cleaned texts before splitting and keep the test partition out of fitting. I report all scores without rounding them up to a target.

Task	Selected model	Accuracy	Macro F1	Weighted F1	Balanced accuracy
Sentiment	LogisticRegression	58.90%	0.5888	0.5873	59.18%
Topic	LogisticRegression	91.73%	0.5920	0.9059	51.71%

My baseline comparison

Task	Baseline accuracy	Baseline macro F1	Selected model errors
sentiment	34.85%	0.1723	487
topic	86.16%	0.2314	98

I prioritize macro F1 because the topic labels are imbalanced. I report weighted F1 and accuracy as supporting measures, not replacements for minority-class recall. I do not claim statistical significance or external generalization from this single split.


My validation comparison

Task	Candidate	Accuracy	Macro F1
sentiment	Majority Baseline	34.77%	0.1720
sentiment	MultinomialNB	57.13%	0.5694
sentiment	LogisticRegression	58.14%	0.5812
sentiment	LinearSVC	56.29%	0.5630
topic	MajorityBaseline	86.08%	0.2313
topic	MultinomialNB	86.16%	0.2360
topic	LogisticRegression	91.14%	0.5546
topic	LinearSVC	91.65%	0.5415

Sentiment evaluation

I report per-class precision, recall, F1, and support below. I read the matrix as actual labels on rows and predicted labels on columns.

Class	Precision	Recall	F1	Support
Negative	0.5950	0.6538	0.6230	364
Neutral	0.5387	0.4891	0.5127	413
Positive	0.6293	0.6324	0.6308	408




I save all 487 errors in `sentiment_errors.csv`. I explain representative mistakes and annotation concerns in my technical report.

Topic evaluation

I report per-class precision, recall, F1, and support below. I read the matrix as actual labels on rows and predicted labels on columns.

Class	Precision	Recall	F1	Support
Account_Security	0.8333	0.2632	0.4000	19
Community_Discussion	0.9239	0.9863	0.9541	1021
Feature_Feedback	0.3684	0.1892	0.2500	37
Technical_Issues	0.9714	0.6296	0.7640	108



I saved all 98 errors in `topic_errors.csv`. I explain representative mistakes and annotation concerns in my technical report.